

Operationalizing an AWS ML Project

- Ting Lu

1. Initial Setup

I have chosen the “**ml.t2.medium**” instance type for Notebook instance (Figure 1 - SageMaker Notebook Instance). There are multiple reasons for selecting this instance type for my notebook.

- Firstly, for completing the execution of this project’s jupyter notebooks we **do not need** a very computationally **powerful CPU** and **high RAM**.
- We will need to keep this **notebook** instance in “**inService**” status for a **long time** while we are working on the project
- In order to **avoid high costs**, we should **select a notebook** that is **low in per hour cost** and offers reasonably **good CPU and RAM**.
- So looking at the instance type and their pricing: <https://aws.amazon.com/SageMaker/pricing/> , we have two choices “ml.t2.medium” and “ml.t3.medium”. Both have 2 vCPU and 4 GB Memory, and as per doc “ml.t3.medium” has a slightly higher cost as it has a fast boot time.
- Now given that we do have a critical requirement for a fast boot time, so we can go ahead with the “**ml.t2.medium**” as it offers the same computational power and is lower in per hour cost.

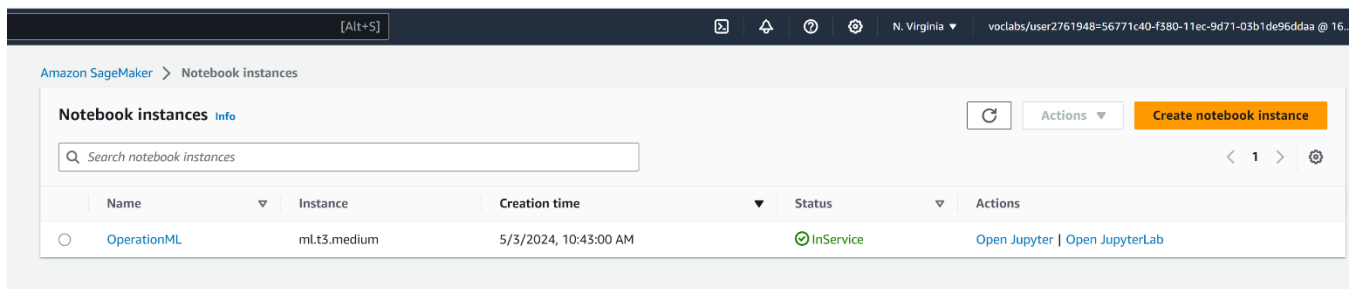


Figure 1. SageMaker Notebook Instance

The dog breed dataset was uploaded to a newly created S3 bucket, successfully.

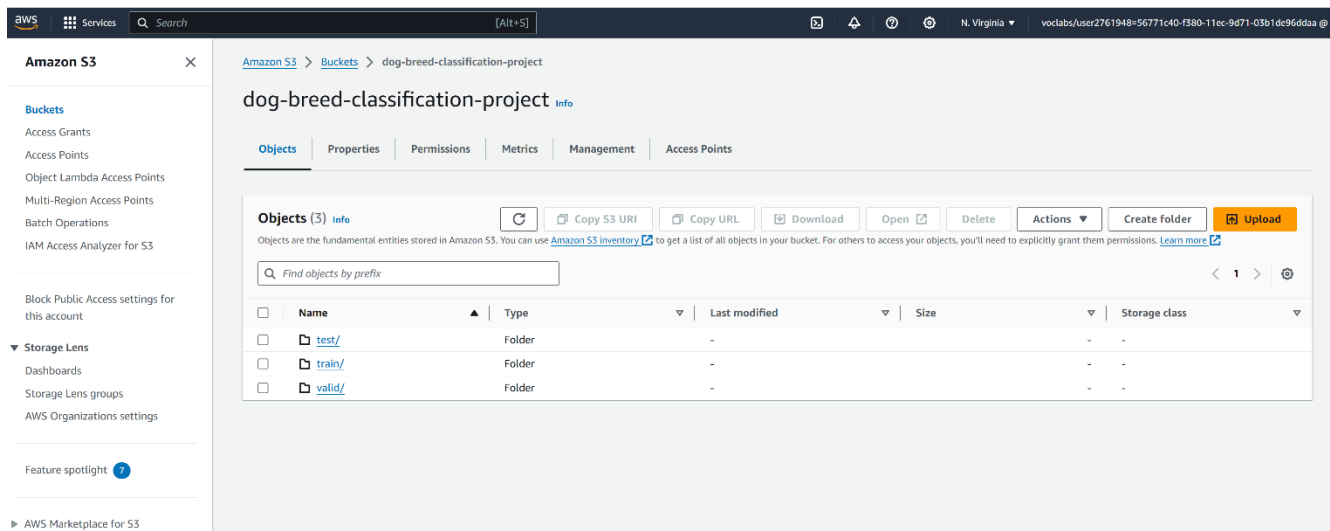


Figure 2. S3 Bucket snapshot

2. SageMaker Training and Deployment

For hyperparameter tuning I used the same “ml.m5.xlarge” instance_type , however since it took too much time for the tuning, while the training process I tried to increase the processing power a bit by using the “ml.m5.2xlarge” for the single instance and multi instance training purposes.

Hyperparameter tuning job (max_jobs = 2, max_parallel_jobs = 2) completed successfully:

Amazon SageMaker > Hyperparameter tuning jobs

Hyperparameter tuning jobs

Search hyperparameter tuning jobs

< 1 >

Name	Status	Training completed/total	Creation time	Duration
pytorch-training-240504-0433	InProgress	0 / 1	5/4/2024, 12:33:33 AM	a few seconds
pytorch-training-240504-0304	Completed	2 / 2	5/3/2024, 11:04:31 PM	an hour
pytorch-training-240504-0044	Completed	2 / 2	5/3/2024, 8:44:41 PM	an hour
pytorch-training-240503-2126	Completed	6 / 6	5/3/2024, 5:26:44 PM	28 minutes
pytorch-training-240503-2044	Completed	6 / 6	5/3/2024, 4:44:49 PM	25 minutes

Figure 3. Hyperparameter Tuning Job

However, upon training the model with the best parameters from above tuning, the model gave a 0 test accuracy! So increased the max_jobs = 6, max_parallel_jobs = 3 and also changed its instance_type = “ml.m5.xlarge” to speed up the computations a bit.

Reran the hyperparameter jobs and it executed successfully:

Hyperparameter tuning job summary

Name
pytorch-training-240504-0044

ARN
arn:aws:sagemaker:us-east-1:166075758371:hyper-parameter-tuning-job/pytorch-training-240504-0044

Status
Completed

Creation time
May 04, 2024 00:44 UTC

Last modified time
May 04, 2024 01:30 UTC

Approx. total training duration
40 minute(s)

Best training job | Training jobs | Training job definitions | Tuning Job configuration | Tags

Training job status counter

Completed 2 | In Progress 0 | Stopped 0 | Failed 0 (Retryable: 0, Non-retryable: 0)

Training jobs

Sorting by objective metric value will display only jobs that have metric values.

Search training jobs

< 1 >

Name	Status	Final objective metric value	Creation time	Training Duration
pytorch-training-240504-0044-002-945318a8	Completed	130	5/3/2024, 9:09:35 PM	19 minute(s)
pytorch-training-240504-0044-001-efc5ae37	Completed	466	5/3/2024, 8:44:46 PM	21 minute(s)

Figure 4. Hyperparameter jobs summary

Post which triggered the **single instance** and **multi-instance training jobs**. Jobs completed successfully. For snapshot of the training jobs refer the images folder in the repository.

Amazon SageMaker > Training jobs > dog-pytorch-2024-05-04-01-30-34-546

dog-pytorch-2024-05-04-01-30-34-546

Clone Create model package Stop Create model

Job settings

Job name dog-pytorch-2024-05-04-01-30-34-546	Status Completed View history	SageMaker metrics time series Enabled	IAM role ARN arn:aws:iam::166075758371:role/service-role/AmazonSageMaker-ExecutionRole-20240503T104235
ARN arn:aws:sagemaker:us-east-1:166075758371:training-job/dog-pytorch-2024-05-04-01-30-34-546	Creation time May 04, 2024 01:30 UTC	Training time (seconds) 734	
	Last modified time May 04, 2024 01:43 UTC	Billable time (seconds) 734	
		Managed spot training savings 0%	
		Tuning job source/parent -	

Algorithm

Algorithm ARN -	Additional volume size (GB) 30	Maximum wait time for managed spot training(s) -	Volume encryption key -
Training image 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.4.0-cpu-py3	Maximum runtime (s) 86400	Managed spot training Disabled	
Input mode File			

Instance group	Instance type	Instance count	Keep alive period
-	m5.xlarge	1	-

Figure 5. Single Instance Training Job

Amazon SageMaker > Training jobs > dog-pytorch-2024-05-04-01-43-56-880

dog-pytorch-2024-05-04-01-43-56-880

Clone Create model package Stop Create model

Job settings

Job name dog-pytorch-2024-05-04-01-43-56-880	Status Completed View history	SageMaker metrics time series Enabled	IAM role ARN arn:aws:iam::166075758371:role/service-role/AmazonSageMaker-ExecutionRole-20240503T104235
ARN arn:aws:sagemaker:us-east-1:166075758371:training-job/dog-pytorch-2024-05-04-01-43-56-880	Creation time May 04, 2024 01:43 UTC	Training time (seconds) 740	
	Last modified time May 04, 2024 01:57 UTC	Billable time (seconds) 740	
		Managed spot training savings 0%	
		Tuning job source/parent -	

Algorithm

Algorithm ARN -	Additional volume size (GB) 30	Maximum wait time for managed spot training(s) -	Volume encryption key -
Training image 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.4.0-cpu-py3	Maximum runtime (s) 86400	Managed spot training Disabled	
Input mode File			

Instance group	Instance type	Instance count	Keep alive period
-	m5.xlarge	5	-

Input data configuration: training

Figure 6. Multi Instance Training Job

Deployed Endpoints:

- Single instance deployed endpoint: “pytorch-inference-2024-05-04-01-57-54-708”
- Multi instance deployed endpoint: “pytorch-inference-2024-05-04-05-46-54-179”

Amazon SageMaker > Endpoints

Endpoints

Search endpoints

Update endpoint Actions Create endpoint

Name	ARN	Creation time	Status	Last updated
pytorch-inference-2024-05-04-01-57-54-708	arn:aws:sagemaker:us-east-1:166075758371:endpoint/pytorch-inference-2024-05-04-01-57-54-708	5/3/2024, 9:57:55 PM	InService	5/3/2024, 10:00:49 PM
pytorch-inference-2024-05-04-05-46-54-179	arn:aws:sagemaker:us-east-1:166075758371:endpoint/pytorch-inference-2024-05-04-05-46-54-179	5/4/2024, 1:46:54 AM	InService	5/4/2024, 1:50:11 AM

Figure 7. SageMaker Endpoints

3. EC2 Training

- We have utilized the **t2.xlarge** instance and the **Deep Learning AMI (Amazon Linux 2) Version 55.0**. This seems like a reasonable balance of performance and affordability.
- As per the documentation, T2 instances can sustain high CPU performance for as long as a workload needs it.
- For most general-purpose workloads, T2 instances will provide ample performance without any additional charges.
- Similarly, because we don't know the duration for which we might need to keep this EC2 instance running for training, it's better to go with a medium size instance so we don't have to pay for a large instance while we're doing setup, debugging and other tasks.

Difference between `ec2train1.py` (EC2 script) and `train_and_deploy-solution.ipynb` + `hpo.py` (SageMaker scripts)

- There is no logic for calling any Estimator or Tuner functions in the EC2 script. The code in the EC2 script is responsible for saving the model to the local path. While in the SageMaker scripts this was handled internally by SageMaker where the model data was stored to a S3 location.
- In the EC2 training script, all the variables like hyperparameters and output locations, etc are already mentioned in the script itself and so there is no need for **argparse**. Meaning while running the EC2 script we do not need to mention any arguments.
- In the EC2 script the training happens on the same server on which the script is invoked/executed, however in the SageMaker scripts the training job that is invoked, it runs on a separate container than the one on which the SageMaker notebook is running.
- Another difference is that `ec2train1.py` lacks the main function
- For the EC2 Training, given that the training data and model, all are stored on the EC2 instance host itself it would be difficult to deploy the saved model to an endpoint in SageMaker. If we wish to do that then we might need to manually upload the model first to SageMaker and then use that to deploy an endpoint. This is not the case in models trained via the SageMaker notebook instances, as the model can be easily deployed to an endpoint.

For snapshot refer below:

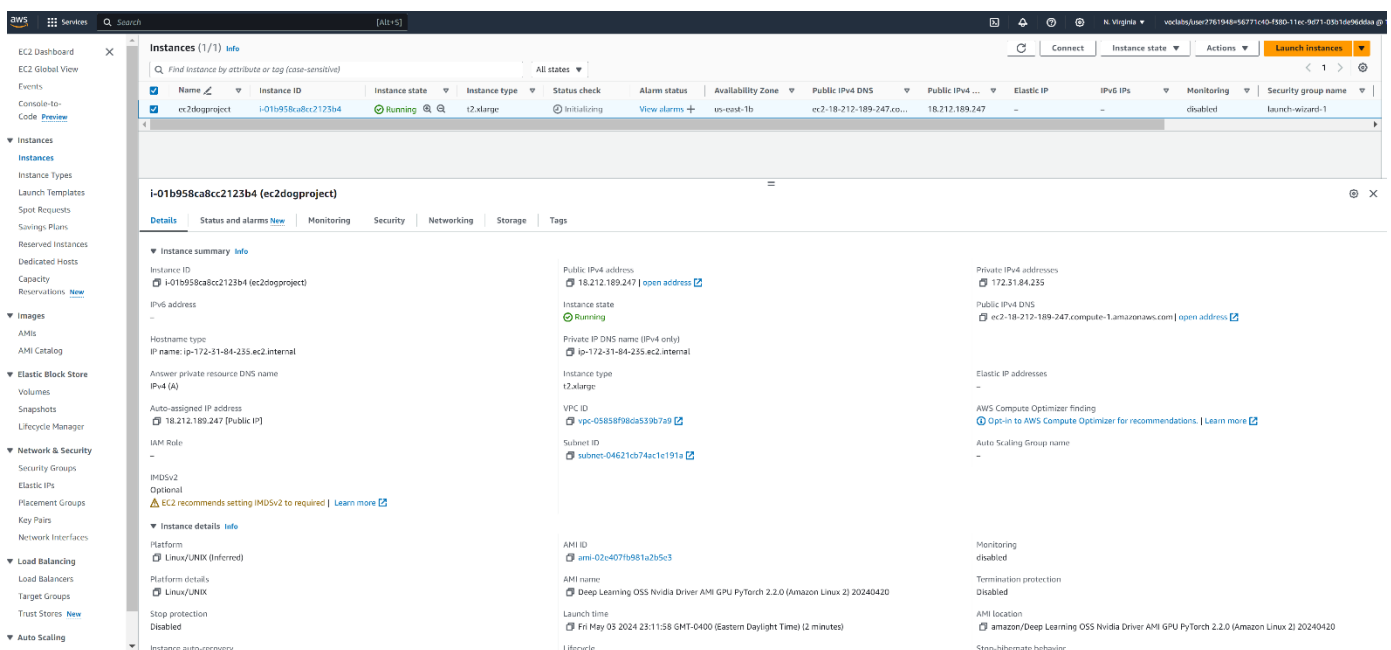


Figure 8. EC2 Instance snapshot

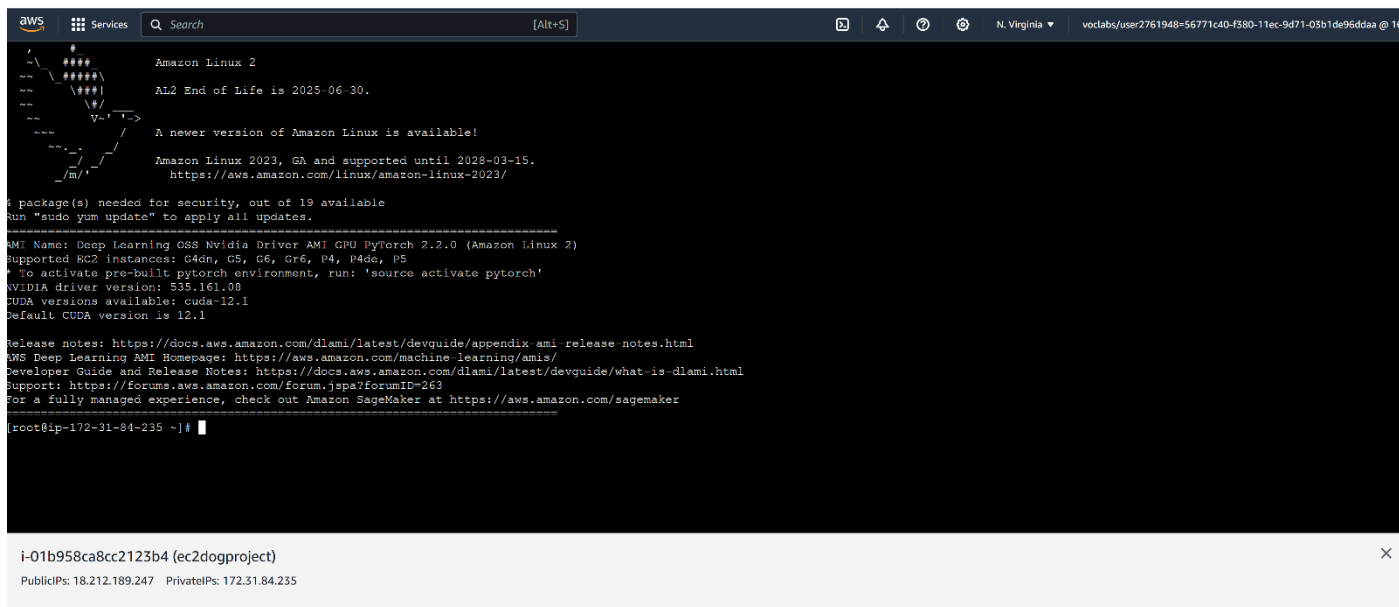


Figure 9. EC2 Terminal

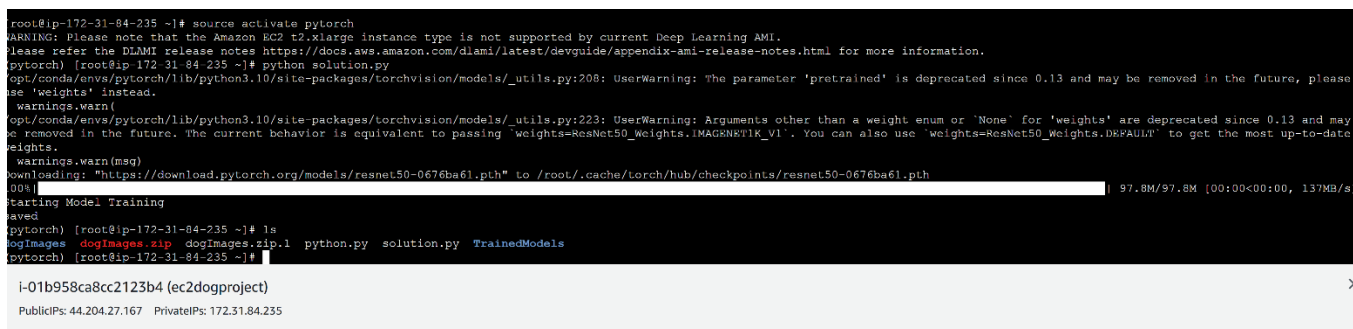


Figure 10. EC2 Training saved - model.pth

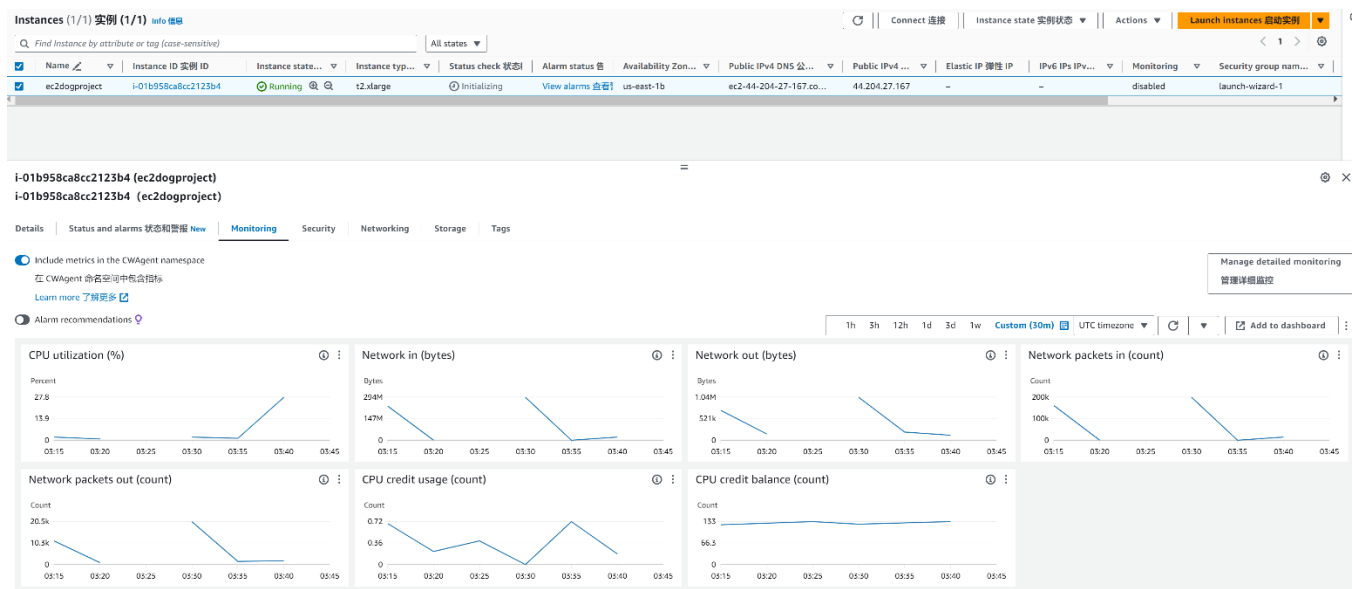


Figure 11. EC2 Instance Resource Metrics

4. Lambda functions

- The lambda functions will be used for invoking our deployed endpoints.
- The lambda function implemented in this project expects the image inputs in json format, which is used to invoke the model's deployed endpoint
- Given we have two endpoints deployed, one for the single instance training and the other for the multi-instance training, we will only use the multi-instance training jobs endpoint and create a lambda function for invoking that endpoint.
- Multi instance trained endpoint that we will be using: **“pytorch-inference-2022-06-08-20-39-47-131”**
- We created the lambda function with the corresponding changes to invoke the endpoint:

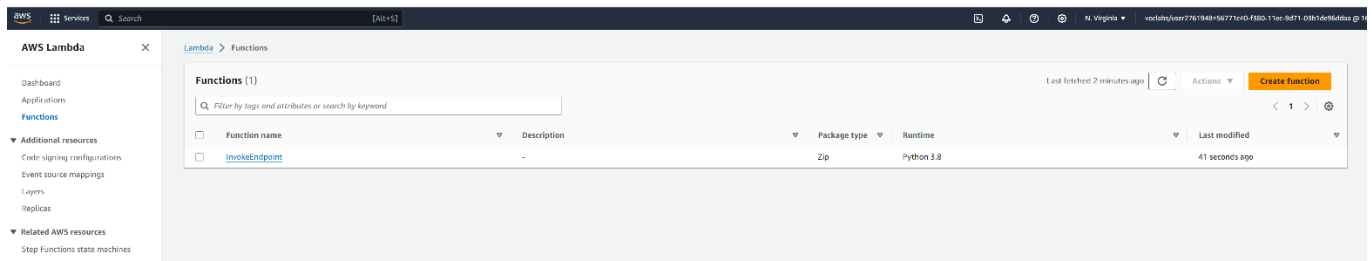


Figure 12. Lambda Function

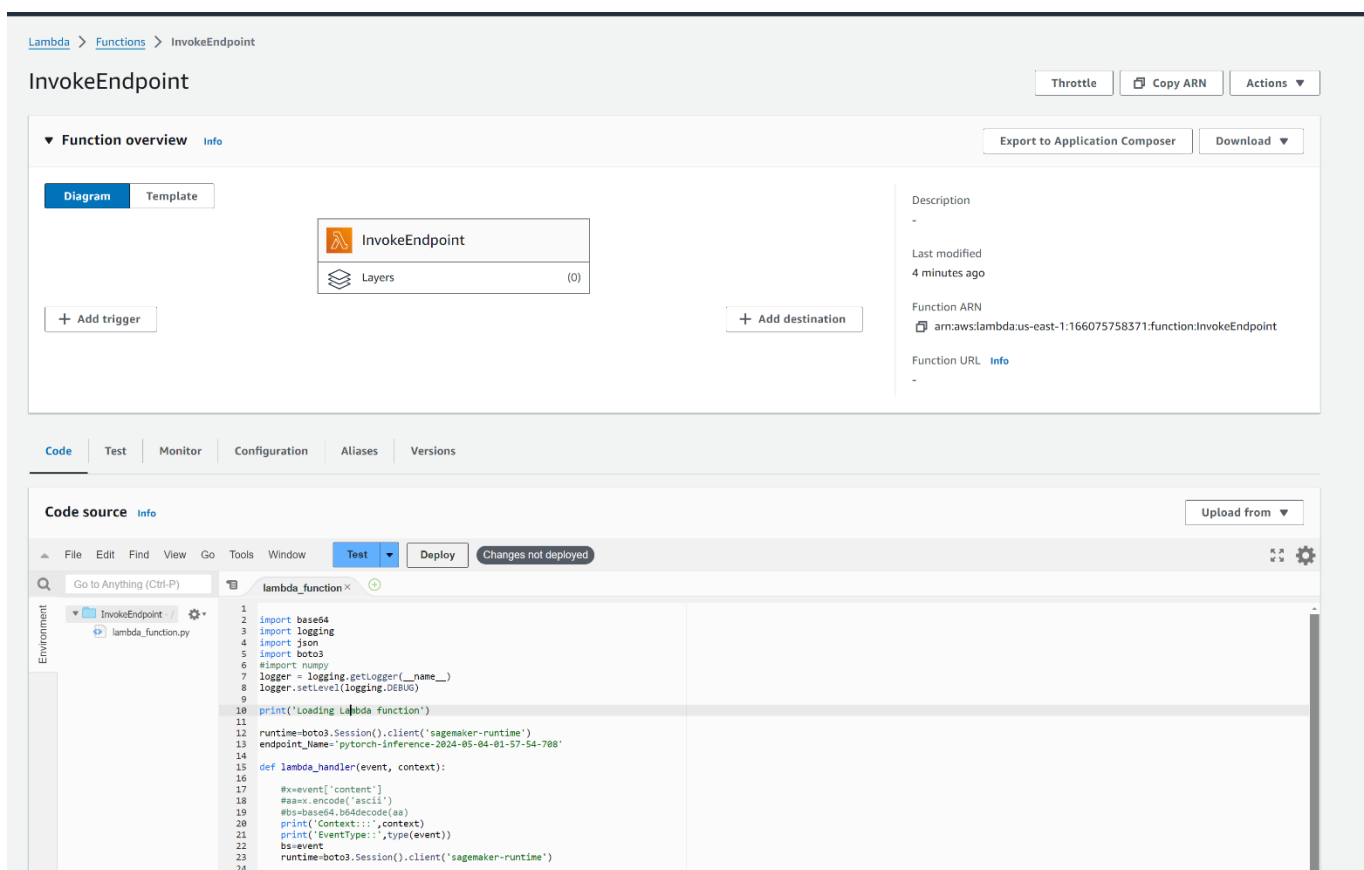


Figure 13. Lambda Function

5. Security and Testing

- We had created a new role for this lambda functions with basic accesses.
- We used the below test event to test our lambda function:

Configure test event

×

A test event is a JSON object that mocks the structure of requests emitted by AWS services to invoke a Lambda function. Use it to see the function's invocation result.

To invoke your function without saving an event, configure the JSON event, then choose Test.

Test event action

Create new event

Edit saved event

Event name

InvokeEndpoint

Maximum of 25 characters consisting of letters, numbers, dots, hyphens and underscores.

Event sharing settings

Private

This event is only available in the Lambda console and to the event creator. You can configure a total of 10. [Learn more](#)

Shareable

This event is available to IAM users within the same account who have permissions to access and use shareable events. [Learn more](#)

Template - optional

hello-world

Event JSON

Format JSON

```

1 {
2   "url": "https://s3.amazonaws.com/cdn-origin-etr.akc.org/wp-content/uploads/2017/11/20113314/Ce
3 }

```

Figure 14. Lambda Function Test Event

- Now when we tried to execute the test event we got and **AccessDeniedException**. This was expected as the lambda function did not have access to invoke the SageMaker endpoint.
- Error Message:

```

{
  "errorMessage": "An error occurred (AccessDeniedException) when calling the InvokeEndpoint operation: User: arn:aws:sts::594529116801:assumed-role/invokeEndpoint-role-4qryi0br/invokeEndpoint is not authorized to perform: sagemaker:InvokeEndpoint on resource: arn:aws:sagemaker:us-east-1:594529116801:endpoint/pytorch-inference-2024-05-03-10-39-47-131 because no identity-based policy allows the sagemaker:InvokeEndpoint action",
  "errorType": "ClientError",
  "stackTrace": [
    " File \"/var/task/lambda_function.py\", line 24, in lambda_handler\n",
    " response=runtime.invoke_endpoint(EndpointName=endpoint_Name,\n",
    " File \"/var/runtime/botocore/client.py\", line 391, in _api_call\n",
    " return self._make_api_call(operation_name, kwargs)\n",
    " File \"/var/runtime/botocore/client.py\", line 719, in _make_api_call\n",
    " raise error_class(parsed_response, operation_name)\n",
  ]
}

```

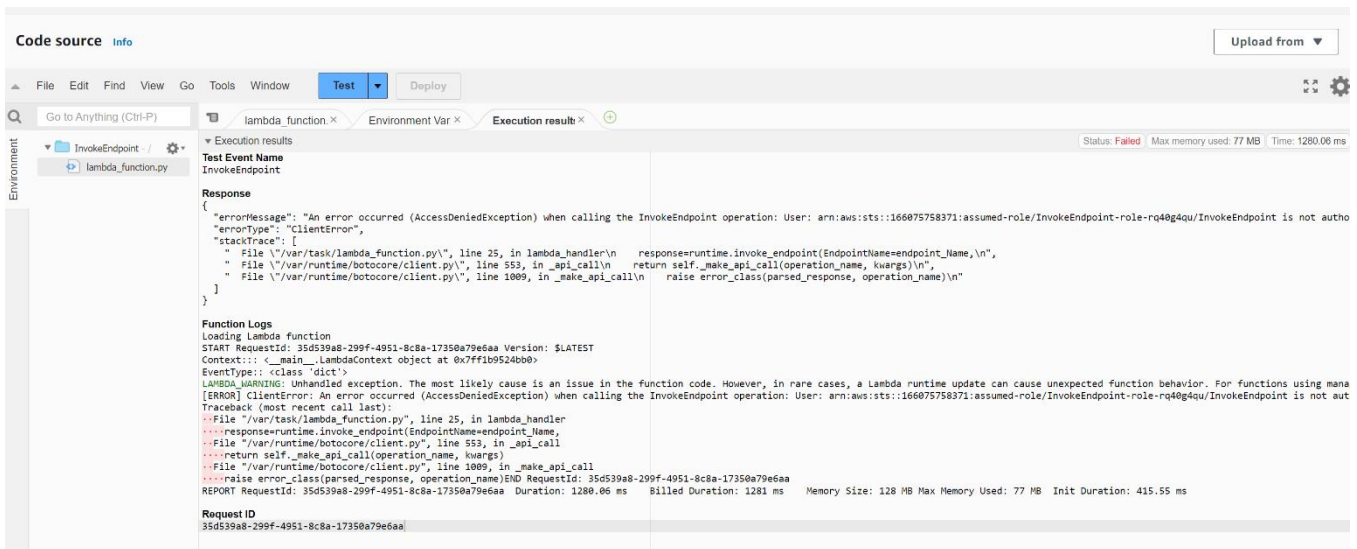


Figure 15. Lambda function test event failure response

- So went ahead and added the “SageMakerFullAccess” policy role to the lambda function’s role.

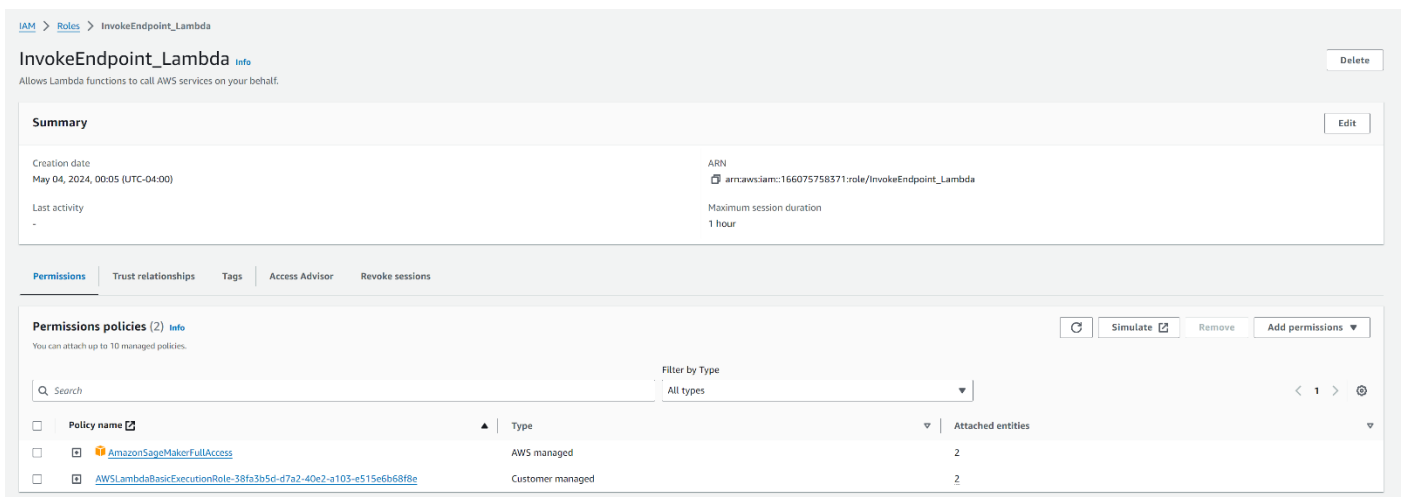


Figure 16. Lambda function role IAM permissions

- Post which we got the following output from the test event:
- (Please note that there are 133 dog breed and not 33 as mentioned in the project instructions. So, we do expect there to be around 133 elements in the response object.)
- Response:

```
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "<__main__.LambdaContext object at 0x7f3802e37b20>",
  "body": "[[-7.732433795928955, -7.176344394683838, -6.378891468048096, -1.5176353454589844, -
3.4210686683654785, -9.79492473602295, -6.717096328735352, -1.7850936651229858, -
8.523921012878418, -1.404198169708252, -0.7228893637657166, -5.840402603149414, -
5.7331461906433105, -2.0368082523345947, -9.905051231384277, -6.480345726013184, -
10.379494667053223, -2.0244674682617188, -10.030835151672363, 0.24102921783924103, -
7.826473712921143, -2.3736279010772705, -7.411016464233398, -9.585257530212402, -
```



```

6.935138702392578, -10.370619773864746, -6.9618120193481445, -5.8000054359436035, -
7.852663516998291, -4.261888027191162, -8.121268272399902, -2.3057727813720703, -
9.783154487609863, -4.807986736297607, -8.92819595336914, -8.51605224609375, -8.291374206542969,
-5.918830394744873, -3.329312324523926, -7.71879768371582, -4.240849018096924, -
8.23236083984375, -0.6575157642364502, -3.2674717903137207, -1.3542909622192383, -
12.321500778198242, -5.892731666564941, -0.7067917585372925, -4.4090352058410645, -
3.0732972621917725, -3.965545892715454, -12.868936538696289, -10.151001930236816, -
5.127017974853516, -9.019100189208984, -2.250539779663086, -6.182336807250977, -
10.275800704956055, -1.599454641342163, -6.101325988769531, -13.15511417388916, -
13.388274192810059, -12.73229694366455, -6.6955180168151855, -3.183396577835083, -
13.372259140014648, -2.6620848178863525, -5.736507892608643, -3.5820565223693848, -
0.9081988334655762, 1.058609127998352, -7.375503063201904, -7.269904136657715, -
5.875174045562744, -8.469480514526367, -4.669737815856934, -8.637510299682617, -
2.2320778369903564, -6.134050369262695, -4.5598530769348145, -2.207737922668457, -
10.790899276733398, -2.3271493911743164, -3.509336471557617, -10.742084503173828, -
11.33560848236084, -7.614548683166504, -12.094978332519531, -5.790018558502197, -
2.1963579654693604, -9.642144203186035, -7.176695346832275, -6.40428352355957, -
11.0121488571167, -6.617325305938721, -3.0795605182647705, -7.987880706787109, -
3.7228918075561523, -9.426275253295898, -7.498056888580322, -11.697783470153809, -
1.4410004615783691, -2.5681569576263428, -4.899290084838867, -3.1299190521240234, -
5.716061115264893, -8.185333251953125, -3.402395248413086, -2.385983467102051, -
2.442678928375244, -8.68392562866211, -2.8638479709625244, -11.222406387329102, -
10.302404403686523, -7.736847877502441, -4.212231636047363, -7.31355619430542, -
4.08905029296875, -12.546906471252441, -2.2557764053344727, -2.3533856868743896, -
6.100186347961426, -7.259090900421143, -6.4662346839904785, -7.800014495849609, -
6.287154197692871, -6.792305946350098, -2.09424090385437, -4.197015762329102, -
12.041462898254395, -10.169670104980469, -1.5505112409591675, -6.9473676681518555]]]"
}

```

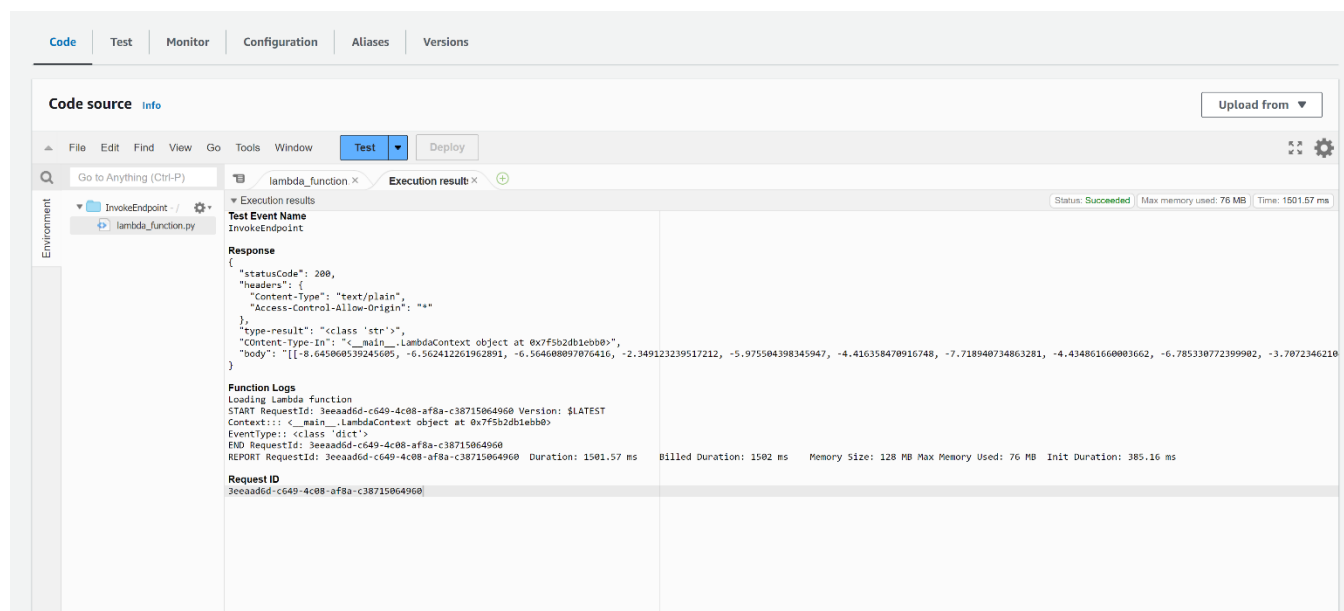


Figure 17. Lambda function test event success response

- I'm concerned about the "Full Access" type permission policies that are available.
- For example, for this lambda functions we have provided the lambda function a Full Access to SageMaker resources, but this does not seem to follow the concept of least privilege access.
- Ideally, one should only allow these lambda functions to query the endpoints that they're intended and allowed to query.

- We will have to do some more analysis to figure out if there's anything we could do about it.
- Furthermore, another concern is that the account's root user does not employ MFA
- Looking at the IAM roles that are currently active, all the roles seem to be necessary and most of the roles have been added on a per need basis.
- However, we need to keep an eye on the roles dashboard to make sure only relevant roles necessary for currently active projects, are the only roles that are in active state to prevent unauthorized accesses.

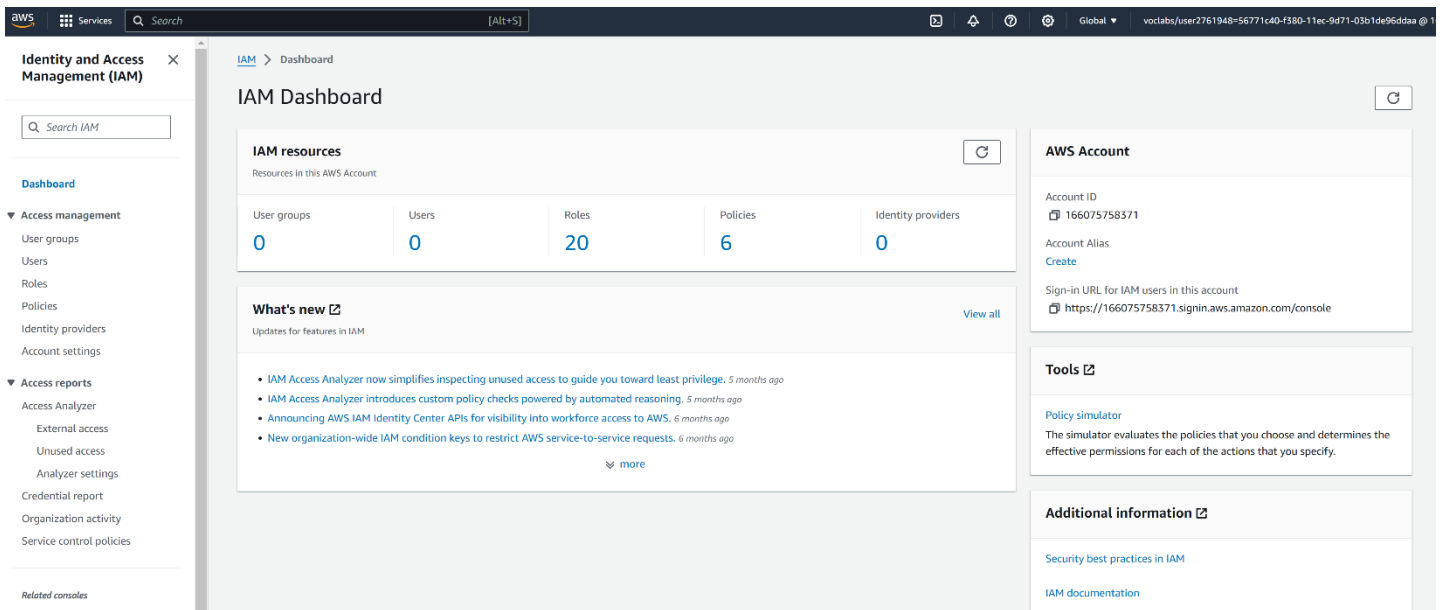


Figure 18. IAM Dashboard

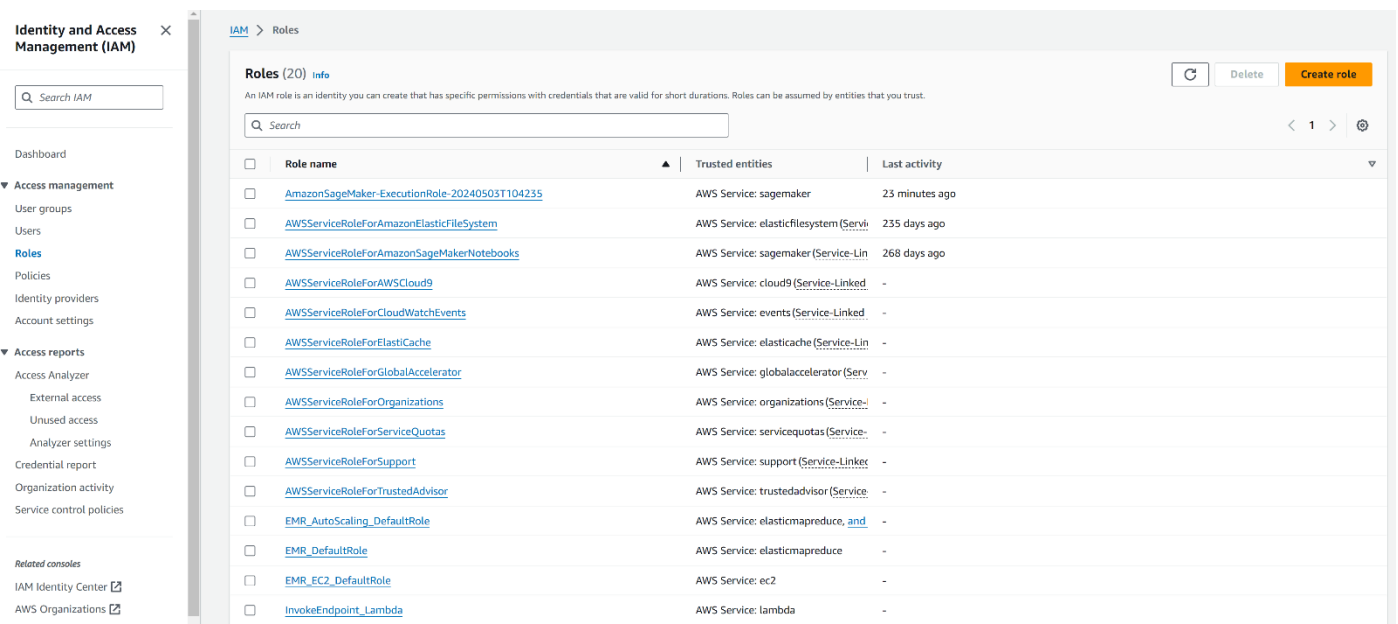


Figure 19. IAM Roles

6. Concurrency and Auto-scaling

- Before adding in configs for Concurrency and Auto-scaling for our lambda functions we will first create a version config for our lambda function.

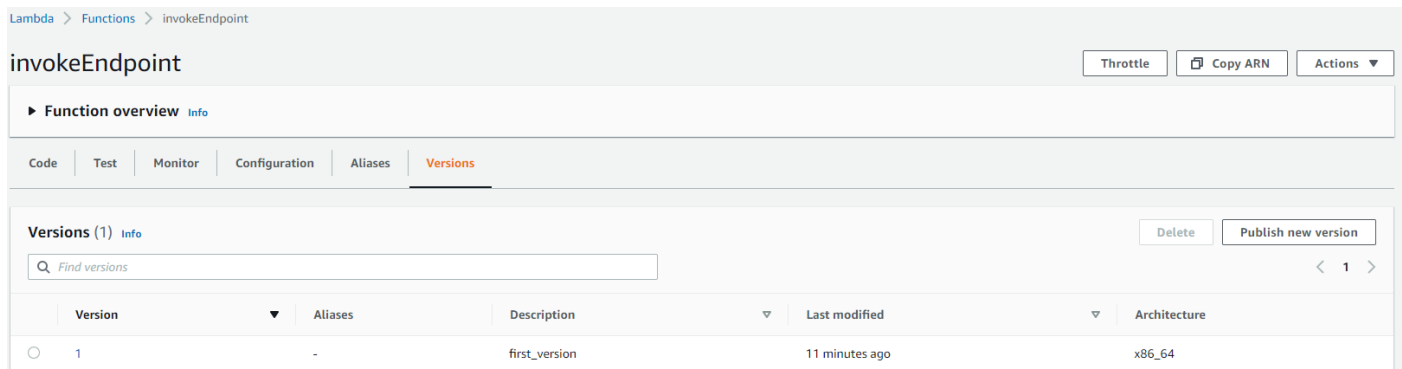


Figure 20. Lambda function version config

- For the lambda function we have set the **reserved concurrency** to be 5. This implies that the lambda function would be able to handle up to 5 requests concurrently at the same time. This would help lower latency issues in situations when there is higher traffic than usual.

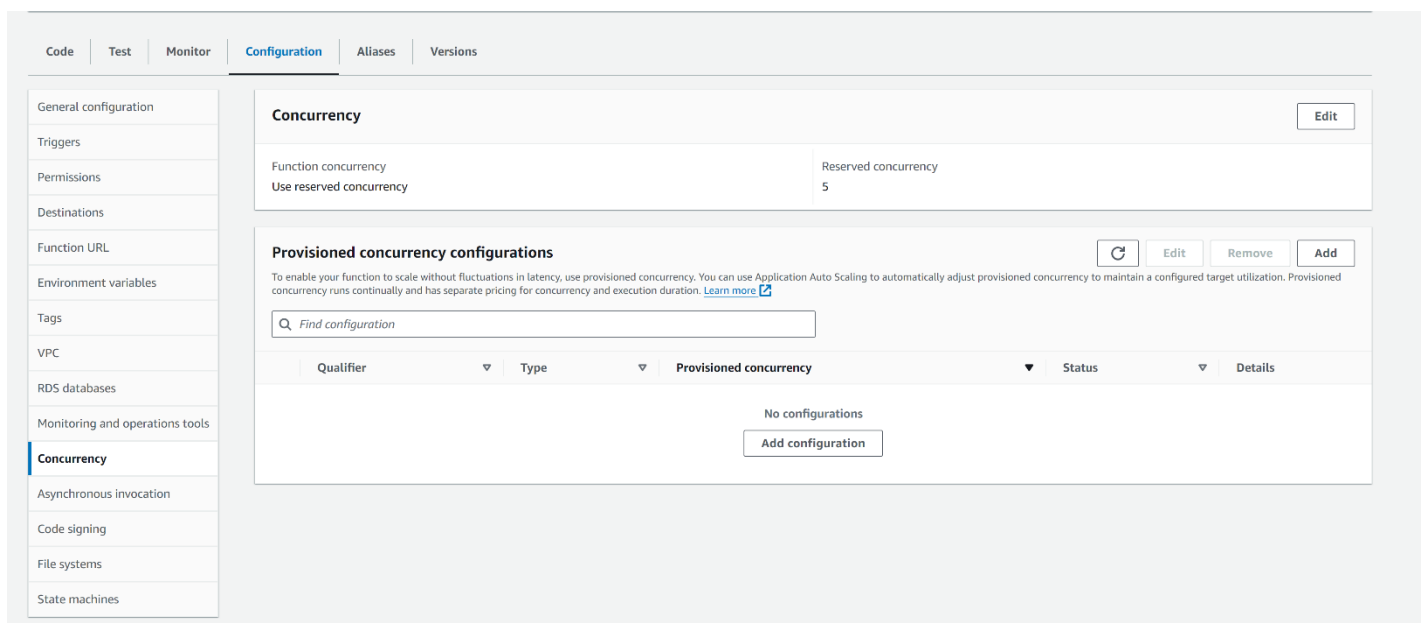


Figure 21. lambda function reserved concurrency

- Given the current use case, ideally using only reserved concurrency should have sufficed for our use case and we might not need to consider using the provisioned concurrency configs. However, for the sake of completion, I tried to add in the config for the provisioned concurrency as well. We set the provision concurrency to be 2.

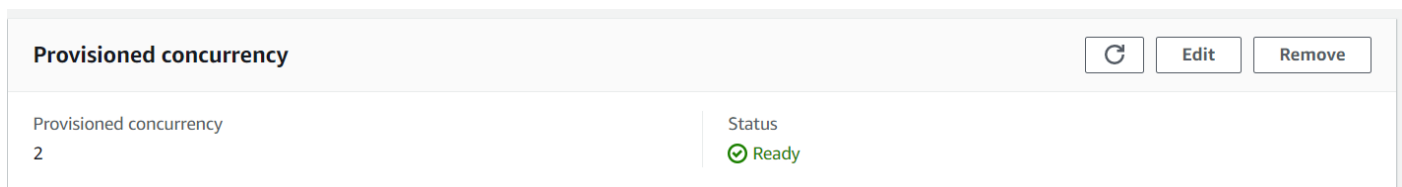


Figure 22. lambda function provisioned concurrency

- For adding config for auto-scaling we have added the below configs:

Amazon SageMaker > Endpoints > pytorch-inference-2024-05-04-01-57-54-708 > AllTraffic

Configure variant automatic scaling

Deregister auto scaling

Variant automatic scaling [Learn more](#)

Variant name
AllTraffic

Instance type
ml.m5.large

Elastic Inference
-

Current instance count
1

Current weight
1

Minimum instance count
1

 -

Maximum instance count
3

IAM role
Amazon SageMaker uses the following service-linked role for automatic scaling. [Learn more](#)

AWSServiceRoleForApplicationAutoScaling_SageMakerEndpoint

Built-in scaling policy [Learn more](#)

Policy name
SageMakerEndpointInvocationScalingPolicy

Target metric
[SageMakerVariantInvocationsPerInstance](#)

Target value
10

Scale in cool down (seconds) - optional
30

Scale out cool down (seconds) - optional
30

☐ Disable scale in

Select if you don't want automatic scaling to delete instances when traffic decreases. [Learn more](#)

Amazon SageMaker > Endpoints > pytorch-inference-2024-05-04-01-57-54-708

pytorch-inference-2024-05-04-01-57-54-708

Delete

Endpoint summary

Name
pytorch-inference-2024-05-04-01-57-54-708

ARN
arn:aws:sagemaker:us-east-1:166075758371:endpoint/pytorch-inference-2024-05-04-01-57-54-708

URL
<https://runtime.sagemaker.us-east-1.amazonaws.com/endpoints/pytorch-inference-2024-05-04-01-57-54-708/invocations>
[Learn more about the API](#)

Status
InService

Creation time
Fri May 03 2024 21:57:55 GMT-0400 (Eastern Daylight Time)

Model container logs
[/aws/sagemaker/endpoints/pytorch-inference-2024-05-04-01-57-54-708](#)

Type
Real-time

Last updated
Fri May 03 2024 22:00:49 GMT-0400 (Eastern Daylight Time)

Alarms
0 alarms

Monitor **Settings** Alarms

Data capture settings

Enable data capture
No

Current sampling percentage (%)
-

S3 location to store data collected
-

Data capture status
-

Endpoint runtime settings

Update weights

Update instance count

Configure auto scaling

	Variant name	Current weight	Desired weight	Elastic Inference	Instance type	Current instance count	Desired instance count	Instance min - max	Automatic scaling
<input type="radio"/>	AllTraffic	1	1	-	ml.m5.large	1	1	1 - 3	Yes

Figure 23. Endpoint Auto-scaling Config

12

- We have set the max instance count to 3 for Auto-scaling, as considering the current requirement, auto scaling on 3 instances with a scale-in and scale-out cool down time of 30 seconds should be a reasonably good config.

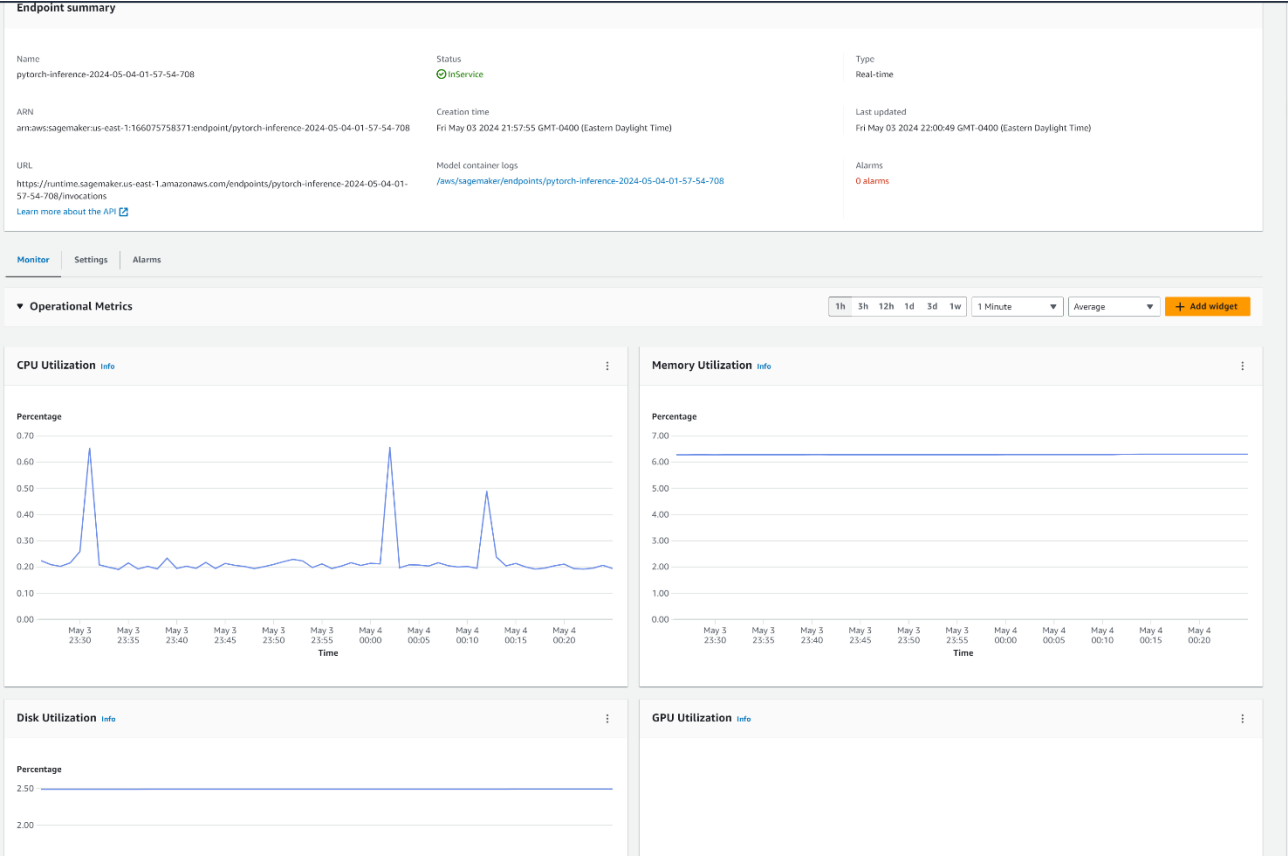


Figure 24. Endpoint Auto-scaling Metrics